

In [32]:

```
import pandas as pd
import numpy as np

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.svm import LinearSVC
from sklearn.naive_bayes import GaussianNB

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, classification_report, confusi

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.utils import resample
from sklearn.preprocessing import StandardScaler

from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report

sns.set_style('whitegrid')
sns.set_palette('pastel')
import warnings

warnings.simplefilter("ignore")
```

In [33]:

```
df = pd.read_csv(r"C:\Users\Simi\Downloads\churn prediction.csv")
```

In [34]:

```
print(df.head())
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
In [35]: print(df.isnull().sum())
```

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

```
In [36]: df.columns
```

```
Out[36]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
               'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
               'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [37]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column             Non-Null Count  Dtype  
---  -
0   RowNumber          10000 non-null  int64  
1   CustomerId         10000 non-null  int64  
2   Surname            10000 non-null  object  
3   CreditScore        10000 non-null  int64  
4   Geography          10000 non-null  object  
5   Gender             10000 non-null  object  
6   Age                10000 non-null  int64  
7   Tenure             10000 non-null  int64  
8   Balance            10000 non-null  float64 
9   NumOfProducts      10000 non-null  int64  
10  HasCrCard          10000 non-null  int64  
11  IsActiveMember     10000 non-null  int64  
12  EstimatedSalary    10000 non-null  float64 
13  Exited             10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [38]: df.head(3)
```

```
Out[38]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Ba
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	838
2	3	15619304	Onio	502	France	Female	42	8	1596



```
In [39]: df.shape
```

```
Out[39]: (10000, 14)
```

```
In [40]: is_Exited = df["Exited"].value_counts()
print("Yes: ",is_Exited[1])
print("No: ",is_Exited[0])
```

```
Yes:  2037
```

```
No:   7963
```

```
In [41]: print(df.isna().sum().sum())
print(df.duplicated().sum())
```

```
0
```

```
0
```

```
In [42]: fig,axb = plt.subplots(ncols=2,nrows=1,figsize=(15, 8))


#Gender Distribution
explode = [0.1, 0.1]
df.groupby('Gender')['Exited'].count().plot.pie(explode=explode, autopct="

ax = sns.countplot(x="Gender", hue="Exited", data=df,ax=axb[1])

# Add values on top of each bar
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.ge
        ha='center', va='center', xytext=(0, 10), textcoords='offs

# Set labels and title
plt.title("Distribution of Gender with Exited Status")
plt.xlabel("Gender")
plt.ylabel("Count")

# Show the plot
plt.show()
```




```

In [43]: is_Exited = df["Exited"].value_counts()
plt.figure(figsize=(10, 5)) # Set the same figsize for both plots
plt.subplot(1, 2, 1) # Subplot for the Exited Counts pie chart
plt.pie(is_Exited, labels=["No", "Yes"], autopct="%0.0f%%")
plt.title("Is Exited Counts")

# Distribution of Geography Pie Chart
plt.subplot(1, 2, 2) # Subplot for the Distribution of Geography pie chart
geography_counts = df['Geography'].value_counts()
plt.pie(geography_counts, labels=geography_counts.index, autopct='%1.1f%%')
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle
plt.title('Distribution of Geography')


plt.tight_layout() # Adjust layout to prevent overlapping
plt.show()

```



```
In [44]: sns.countplot(x='Geography',hue='Exited',data=df, palette='rainbow')
```

```
Out[44]: <Axes: xlabel='Geography', ylabel='count'>
```



```
In [45]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1)

df['Balance'] = df['Balance'].astype(int)
df['EstimatedSalary'] = df['EstimatedSalary'].astype(int)
```

```
In [46]: df.head(3)
```

```
Out[46]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is/
0	619	France	Female	42	2	0	1	1	
1	608	Spain	Female	41	1	83807	1	0	
2	502	France	Female	42	8	159660	3	1	

```
In [47]: # Initialize label encoders
le = LabelEncoder()
# Fit and transform the data
df['Gender'] = le.fit_transform(df['Gender'])
df['Geography'] = le.fit_transform(df['Geography'])

df.head(4)
```

```
Out[47]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is/
0	619	0	0	42	2	0	1	1	
1	608	2	0	41	1	83807	1	0	
2	502	0	0	42	8	159660	3	1	
3	699	0	0	39	1	0	2	0	

```
In [48]: No_class = df[df["Exited"]==0]
yes_class = df[df["Exited"]==1]

No_class = resample(No_class, replace=False, n_samples=len(yes_class))
down_samples = pd.concat([yes_class, No_class], axis=0)

X = down_samples.drop("Exited", axis=1)
y = down_samples["Exited"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```

In [49]: # Count the occurrences of each class in the original dataset
original_class_counts = df["Exited"].value_counts()

# Count the occurrences of each class in the downsampled dataset
downsampled_class_counts = down_samples["Exited"].value_counts()

# Calculate the percentage of each class
original_percentages = original_class_counts / len(df) * 100
downsampled_percentages = downsampled_class_counts / len(down_samples) * 100

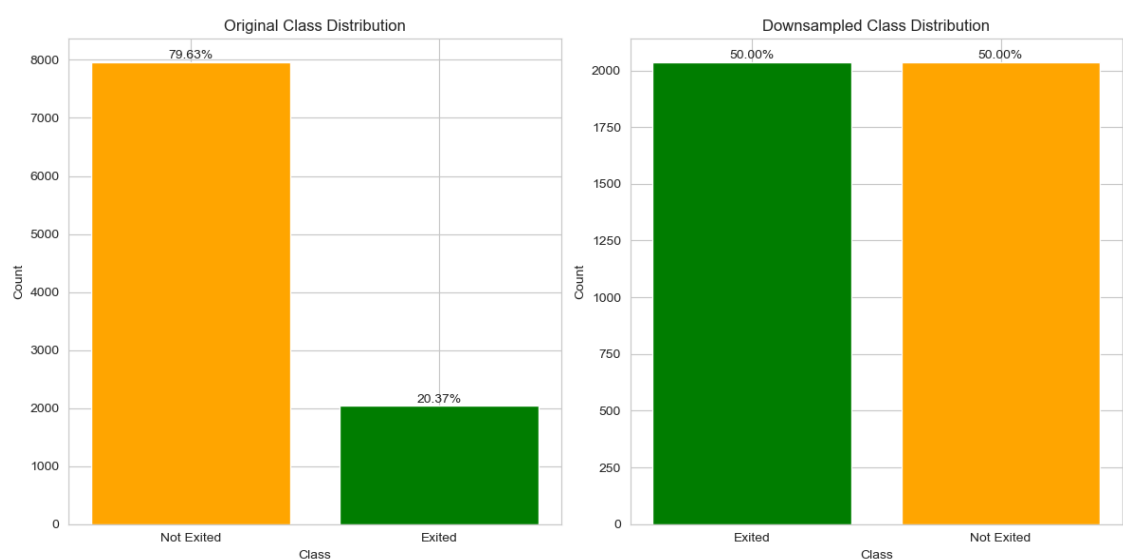
# Plotting
plt.figure(figsize=(12, 6))

# Bar chart for original class distribution
plt.subplot(1, 2, 1)
bars_1 = plt.bar(original_class_counts.index, original_class_counts.values)
for bar, label in zip(bars_1, original_percentages):
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() + 5, f'{label}%')
plt.title('Original Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(original_class_counts.index, ['Not Exited', 'Exited'])

# Bar chart for downsampled class distribution
plt.subplot(1, 2, 2)
bars_2 = plt.bar(downsampled_class_counts.index, downsampled_class_counts.values)
for bar, label in zip(bars_2, downsampled_percentages):
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() + 5, f'{label}%')
plt.title('Downsampled Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(downsampled_class_counts.index, ['Not Exited', 'Exited'])

plt.tight_layout() # the plots will be automatically adjusted to ensure th
plt.show()

```





```
In [50]: DT = DecisionTreeClassifier(max_depth=(5), random_state=0)
DT.fit(X_train, y_train)
predict_ID3 = DT.predict(X_test)
print(classification_report(y_test, predict_ID3))
ID3_accuracy = accuracy_score(predict_ID3,y_test)
print('ID3 model accuracy is: {:.2f}%'.format(ID3_accuracy*100))
```

	precision	recall	f1-score	support
0	0.76	0.82	0.79	410
1	0.80	0.73	0.76	405
accuracy			0.78	815
macro avg	0.78	0.78	0.78	815
weighted avg	0.78	0.78	0.78	815

ID3 model accuracy is: 77.67%

```
In [51]: LR_model = LogisticRegression()
LR_model.fit(X_train, y_train)
predict_LR = LR_model.predict(X_test)
print(classification_report(y_test, predict_LR))
LR_accuracy = accuracy_score(predict_LR,y_test)
print('Logistic Regression accuracy is: {:.2f}%'.format(LR_accuracy*100))
```

	precision	recall	f1-score	support
0	0.72	0.71	0.71	410
1	0.71	0.71	0.71	405
accuracy			0.71	815
macro avg	0.71	0.71	0.71	815
weighted avg	0.71	0.71	0.71	815

Logistic Regression accuracy is: 71.29%

```
In [52]: svm_model = LinearSVC()
svm_model.fit(X_train, y_train)
predict = svm_model.predict(X_test)

print(classification_report(y_test, predict))
svm_accuracy = accuracy_score(predict,y_test)
print('SVC model accuracy is: {:.2f}%'.format(svm_accuracy*100))
```

	precision	recall	f1-score	support
0	0.71	0.72	0.72	410
1	0.71	0.71	0.71	405
accuracy			0.71	815
macro avg	0.71	0.71	0.71	815
weighted avg	0.71	0.71	0.71	815

SVC model accuracy is: 71.29%

```
In [53]: RF = RandomForestClassifier(n_estimators=60, random_state=0)
RF.fit(X_train, y_train)

predict_RF = RF.predict(X_test)

# Evaluate the model
print(classification_report(y_test, predict_RF))
RF_accuracy = accuracy_score(predict_RF, y_test)
print('Random Forest model accuracy is: {:.2f}%'.format(RF_accuracy * 100))
```

	precision	recall	f1-score	support
0	0.77	0.82	0.79	410
1	0.81	0.75	0.77	405
accuracy			0.78	815
macro avg	0.79	0.78	0.78	815
weighted avg	0.79	0.78	0.78	815

Random Forest model accuracy is: 78.40%

```
In [54]: gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate

gb_classifier.fit(X_train, y_train)
y_pred = gb_classifier.predict(X_test)

# Generate classification report
report = classification_report(y_test, y_pred)
print("Classification Report:\n", report)

# Calculate accuracy
gb_accuracy = accuracy_score(y_test, y_pred)
print('XGBoost model accuracy is: {:.2f}%'.format(gb_accuracy * 100))
```

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.83	0.81	410
1	0.82	0.78	0.80	405
accuracy			0.80	815
macro avg	0.80	0.80	0.80	815
weighted avg	0.80	0.80	0.80	815

XGBoost model accuracy is: 80.37%

In [ ]: