

```
In [1]: import numpy as np
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, pre
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.preprocessing import LabelEncoder
```

```
In [4]: df = pd.read_csv(r"C:\Users\Simi\Downloads\email_spam.csv",encoding='latin-1', engine =
```

```
In [5]: df.head()
```

Out[5]:

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [6]: df.tail()
```

Out[6]:

	v1	v2
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will i_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

```
In [7]: df.shape
```

Out[7]: (5572, 2)

```
In [8]: df.size
```

Out[8]: 11144

```
In [9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    v1      5572 non-null    object
 1    v2      5572 non-null    object
dtypes: object(2)
memory usage: 87.2+ KB
```

```
In [10]: df.columns = ["label", "message"]
```

```
In [11]: df.head()
```

```
Out[11]:
```

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [12]: df.isnull().values.any()
```

```
Out[12]: False
```

```
In [13]: df.isnull().sum()
```

```
Out[13]: label      0
message      0
dtype: int64
```

```
In [14]: df.duplicated().values.any()
```

```
Out[14]: True
```

```
In [15]: df.duplicated().sum()
```

```
Out[15]: 403
```

```
In [16]: df.drop_duplicates(inplace=True)
```

```
In [17]: df.describe()
```


```
Out[17]:
```

	label	message
count	5169	5169
unique	2	5169
top	ham	Go until jurong point, crazy.. Available only ...
freq	4516	1

```
In [18]: sns.displot(df.label, kde = True, color = "red")
```

C:\Users\Simi\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)

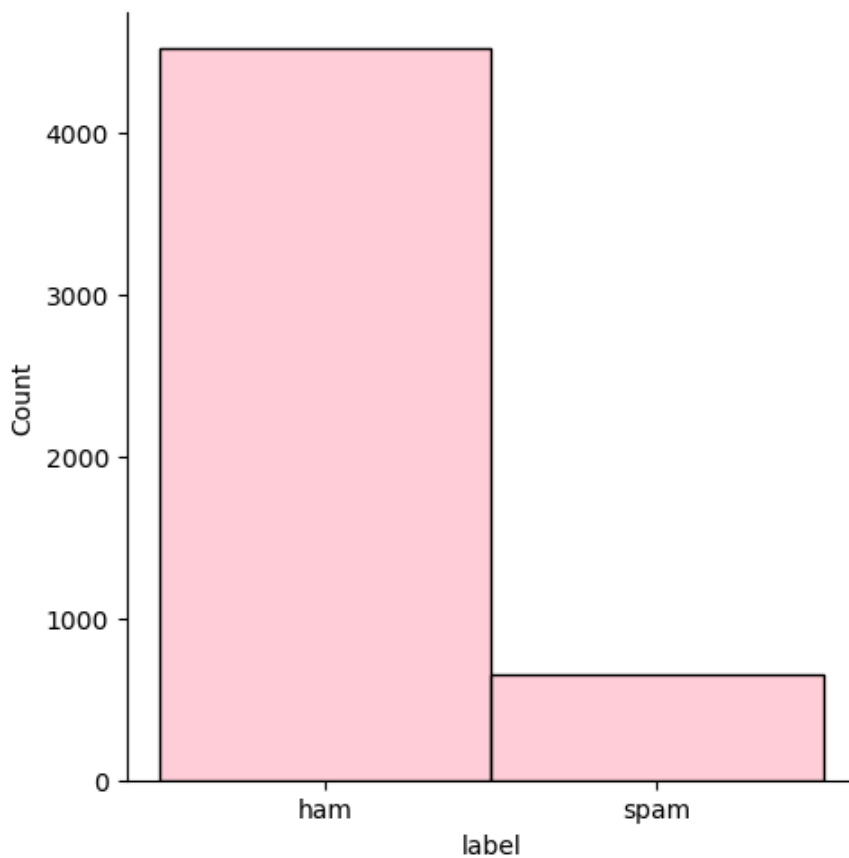
```
Out[18]: <seaborn.axisgrid.FacetGrid at 0x215bf03d590>
```



```
In [19]: sns.displot(df.label, color = "pink")
```

C:\Users\Simi\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)

```
Out[19]: <seaborn.axisgrid.FacetGrid at 0x215bf0e80d0>
```



```
In [20]: encoder = LabelEncoder()  
df["label"] = encoder.fit_transform(df["label"].values)
```

```
In [21]: vectorizer = TfidfVectorizer()
```


```
In [22]: X = vectorizer.fit_transform(df["message"])  
X.toarray()
```

```
Out[22]: array([[0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                ...,  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.]])
```

```
In [23]: y = df["label"]
```

```
In [24]: p, k = len(df[df["label"] == 0]), len(df[df["label"] == 1])
```

```
In [25]: label = np.array(["Ham", "Spam"])
values = np.array([p, k])
plt.figure(figsize=(5, 5))
plt.pie(values, labels=label, autopct="%.1f%%")
plt.show()
```



```
In [26]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [27]: naive_bayes_model = MultinomialNB()
```

```
In [28]: naive_bayes_model.fit(X_train, y_train)
```

```
Out[28]: ▾ MultinomialNB
MultinomialNB()
```

```
In [29]: nb_predictions = naive_bayes_model.predict(X_test)
```

```
In [30]: print("Naive Bayes Model:")
print(confusion_matrix(y_test, nb_predictions))
print(classification_report(y_test, nb_predictions))
print("Accuracy: ", accuracy_score(y_test, nb_predictions))
print("r2_Score: ", r2_score(y_test, nb_predictions))
print("Precision_score: ", precision_score(y_test, nb_predictions))
print("Recall_score: ", recall_score(y_test, nb_predictions))
print("f1_score: ", f1_score(y_test, nb_predictions))
```

Naive Bayes Model:

```
[[889  0]
 [ 46 99]]
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	889
1	1.00	0.68	0.81	145
accuracy			0.96	1034
macro avg	0.98	0.84	0.89	1034
weighted avg	0.96	0.96	0.95	1034

```
Accuracy: 0.9555125725338491
r2_Score: 0.6310150886311625
Precision_score: 1.0
Recall_score: 0.6827586206896552
f1_score: 0.8114754098360656
```

```
In [31]: logistic_regression_model = LogisticRegression()
```

```
In [32]: logistic_regression_model.fit(X_train, y_train)
```

```
Out[32]: ▾ LogisticRegression
LogisticRegression()
```

```
In [33]: lr_predictions = logistic_regression_model.predict(X_test)
```

```
In [34]: print("Logistic Regression Model:")
print(confusion_matrix(y_test, lr_predictions))
print(classification_report(y_test, lr_predictions))
print("Accuracy: ", accuracy_score(y_test, lr_predictions))
print("r2_Score: ", r2_score(y_test, lr_predictions))
print("Precision_score: ", precision_score(y_test, lr_predictions))
print("Recall_score: ", recall_score(y_test, lr_predictions))
print("f1_score: ", f1_score(y_test, lr_predictions))
```

Logistic Regression Model:

```
[[886  3]
 [ 43 102]]
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	889
1	0.97	0.70	0.82	145
accuracy			0.96	1034
macro avg	0.96	0.85	0.90	1034
weighted avg	0.96	0.96	0.95	1034

```
Accuracy: 0.9555125725338491
r2_Score: 0.6310150886311625
Precision_score: 0.9714285714285714
Recall_score: 0.7034482758620689
f1_score: 0.8160000000000001
```

```
In [35]: svm_model = SVC(kernel='linear')
```

```
In [36]: svm_model.fit(X_train, y_train)
```

```
Out[36]: SVC
SVC(kernel='linear')
```

```
In [37]: svm_predictions = svm_model.predict(X_test)
```

```
In [38]: print("Support Vector Machine (SVM) Model:")
print(confusion_matrix(y_test, svm_predictions))
print(classification_report(y_test, svm_predictions))
print("Accuracy: ", accuracy_score(y_test, svm_predictions))
print("r2_Score: ", r2_score(y_test, svm_predictions))
print("Precision_score: ", precision_score(y_test, svm_predictions))
print("Recall_score: ", recall_score(y_test, svm_predictions))
print("f1_score: ", f1_score(y_test, svm_predictions))
```

Support Vector Machine (SVM) Model:

```
[[886  3]
 [ 14 131]]
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	889
1	0.98	0.90	0.94	145
accuracy			0.98	1034
macro avg	0.98	0.95	0.96	1034
weighted avg	0.98	0.98	0.98	1034

```
Accuracy: 0.9835589941972921
r2_Score: 0.8636360110158644
Precision_score: 0.9776119402985075
Recall_score: 0.903448275862069
f1_score: 0.939068100358423
```

In [ ]: